

# Research Brief: federated learning for medical imaging

*Run 20260712-004323 · 3 source(s) summarised.*

**Research question:** what are the privacy risks?

## Introduction

Federated learning (FL) has emerged as a promising approach to train machine learning models while preserving data privacy by decentralizing training across clients. However, as FL systems increasingly adopt medical imaging applications, understanding the specific privacy risks associated with this domain becomes critical. This brief examines the privacy risks in federated learning for medical imaging, focusing on the research question of what constitutes a privacy risk in this context. The brief synthesizes insights from recent literature to provide a clear picture of the current state of knowledge and gaps in understanding.

## Summary

The literature identifies federated learning as a method that enhances privacy through decentralized data processing, but the extent of privacy risks in medical imaging applications remains underexplored. [1] demonstrates that VAFL, a vertical asynchronous federated learning method, employs perturbed local embedding to ensure privacy and communication efficiency in medical imaging contexts, while theoretically analyzing convergence and privacy guarantees across multiple optimization scenarios. [2] highlights vulnerabilities in federated learning systems, particularly when transfer learning is integrated, as this can introduce security risks that compromise privacy. [3] emphasizes FL's capacity to mitigate systemic privacy risks compared to centralized approaches but also underscores the open challenges in ensuring robust privacy protection across diverse client environments.

## Findings

### 1. VAFL: a Method of Vertical Asynchronous Federated Learning

*Tianyi Chen, Xiao Jin, Yuejiao Sun, Wotao Yin (2020)*

*Source reputation: medium — Open-access preprint server; dominant venue for ML/physics/CS research but not peer-reviewed*

*Based on the abstract only — the full text was not accessible.*

The paper introduces VAFL, a vertical asynchronous federated learning method that enables clients to run stochastic gradient algorithms independently without coordination, ensuring suitability for intermittent connectivity. The method employs perturbed local embedding for privacy and communication efficiency while theoretically analyzing convergence rates and privacy levels for strongly convex, nonconvex, and nonsmooth objectives. Empirical results on image and healthcare datasets show VAFL outperforms centralized and synchronous FL methods.

### Key claims:

- VAFL allows each client to run stochastic gradient algorithms without coordination with other clients

The new method allows each client to run stochastic gradient algorithms without coordination with other clients

- VA using perturbed local embedding ensures data privacy and improves communication efficiency

This method further uses a new technique of perturbed local embedding to ensure data privacy and improve communication efficiency

- VAFL theoretically analyzes convergence rates and privacy levels for strongly convex, nonconvex, and nonsmooth objectives

Theoretically, we present the convergence rate and privacy level of our method for strongly convex, nonconvex and even nonsmooth objectives separately

- Empirical results on image and healthcare datasets show VAFL outperforms centralized and synchronous FL methods

Empirically, we apply our method to FL on various image and healthcare datasets. The results compare favorably to centralized and synchronous FL methods

## 2. Federated and Transfer Learning: A Survey on Adversaries and Defense Mechanisms

*Ehsan Hallaji, Roozbeh Razavi-Far, Mehrdad Saif (2022)*

*Source reputation: medium — Open-access preprint server; dominant venue for ML/physics/CS research but not peer-reviewed*

*Based on the abstract only — the full text was not accessible.*

The abstract describes a survey on federated and transfer learning from a security perspective, highlighting how transfer learning addresses security constraints in federated learning. It aims to identify vulnerabilities and defense mechanisms for privacy and performance in such systems.

### Key claims:

- The advent of federated learning has facilitated large-scale data exchange amongst machine learning models while maintaining privacy.

The advent of federated learning has facilitated large-scale data exchange amongst machine: learning models while maintaining privacy.

- One of the most significant advancements in this domain is the incorporation of transfer learning into federated learning, which overcomes fundamental constraints of primary federated learning, particularly in terms of security.

One of the most significant advancements in this domain is the incorporation of transfer learning into federated learning, which overcomes fundamental constraints of primary federated learning, particularly in terms of security.

- This chapter performs a comprehensive survey on the intersection of federated and transfer learning from a security point of view.

This chapter performs a comprehensive survey on the intersection of federated and transfer learning from a security point of view.

- The main goal of this study is to uncover potential vulnerabilities and defense mechanisms that might compromise the privacy and performance of systems that use federated and transfer learning.

The main goal of this study is to uncover potential vulnerabilities and defense mechanisms that might compromise the privacy and performance of systems that use federated and transfer learning.

### 3. Advances and Open Problems in Federated Learning

*Peter Kairouz, H. Brendan McMahan (2020)*

*Source reputation: high — Foundations and Trends® in Machine Learning is a peer-reviewed journal series published by now publishers, with rigorous editorial standards and high impact in the machine learning field.*

*Based on the abstract only — the full text was not accessible.*

This monograph provides an overview of recent advances and open challenges in federated learning (FL), highlighting its decentralized data training approach that enhances privacy while maintaining collaborative model development across clients. The text emphasizes FL's role in addressing systemic privacy risks and costs compared to centralized machine learning, and notes the monograph's focus on summarizing current research progress and identifying unresolved issues in the field.

#### Key claims:

- Federated learning (FL) is a machine learning setting where many clients collaboratively train a model under the orchestration of a central server while keeping training data decentralized.

Federated learning (FL) is a machine as well as a machine learning setting where many clients (e.g., mobile devices or whole organizations) collaboratively train a model under the orchestration of a central server (e.g., service provider), while keeping the training data decentralized.

- FL embodies the principles of focused data collection and minimization, and can mitigate many systemic privacy risks and costs resulting from traditional, centralized machine learning and data science approaches.

FL embodies the principles of focused data collection and minimization, and can mitigate many of the systemic privacy risks and costs resulting from traditional, centralized machine learning and data science

approaches.

- This monograph discusses recent advances and presents an extensive collection of open problems and challenges.

Motivated by the explosive growth in FL research, this monograph discusses recent advances and presents an extensive collection of open problems and challenges.

## Claims and hypotheses

- VAFL allows each client to run stochastic gradient algorithms without coordination with other clients [1]
- VA using perturbed local embedding ensures data privacy and improves communication efficiency [1]
- VAFL theoretically analyzes convergence rates and privacy levels for strongly convex, nonconvex, and nonsmooth objectives [1]
- Empirical results on image and healthcare datasets show VAFL outperforms centralized and synchronous FL methods [1]
- The advent of federated learning has facilitated large-scale data exchange amongst machine learning models while maintaining privacy. [2]
- One of the most significant advancements in this domain is the incorporation of transfer learning into federated learning, which overcomes fundamental constraints of primary federated learning, particularly in terms of security. [2]
- This chapter performs a comprehensive survey on the intersection of federated and transfer learning from a security point of view. [2]
- The main goal of this study is to uncover potential vulnerabilities and defense mechanisms that might compromise the privacy and performance of systems that use federated and transfer learning. [2]
- Federated learning (FL) is a machine learning setting where many clients collaboratively train a model under the orchestration of a central server while keeping training data decentralized. [3]
- FL embodies the principles of focused data collection and minimization, and can mitigate many systemic privacy risks and costs resulting from traditional, centralized machine learning and data science approaches. [3]
- This monograph discusses recent advances and presents an extensive collection of open problems and challenges. [3]

## Conclusion

The evidence indicates that federated learning for medical imaging presents significant privacy risks, particularly when systems involve asynchronous communication or transfer learning integration, though VAFL's perturbed local embedding approach demonstrates a mechanism for enhancing privacy while maintaining model performance.

## Further research

- How does the perturbed local embedding in VAFL specifically address privacy risks in medical imaging data without compromising diagnostic accuracy?
- What are the security vulnerabilities introduced by transfer learning in federated medical imaging systems?
- To what extent do the privacy guarantees of VAFL hold in real-world medical imaging

environments with heterogeneous client data distributions?

- How do asynchronous communication patterns in federated learning systems impact the robustness of privacy mechanisms in medical imaging applications?

## References

- Tianyi Chen, Xiao Jin, Yuejiao Sun, Wotao Yin (2020) 'VAFL: a Method of Vertical Asynchronous Federated Learning', arXiv:2007.06081v1. Available at: <http://arxiv.org/abs/2007.06081v1> [reputation: medium] [access: open]
- Ehsan Hallaji, Roozbeh Razavi-Far, Mehrdad Saif (2022) 'Federated and Transfer Learning: A Survey on Adversaries and Defense Mechanisms', arXiv:2207.02337v1. Available at: <http://arxiv.org/abs/2207.02337v1> [reputation: medium] [access: open]
- Peter Kairouz, H. Brendan McMahan (2020) 'Advances and Open Problems in Federated Learning', doi:10.1561/22000000083. Available at: <https://openalex.org/W2995022099> [reputation: high] [citations: 4812] [access: open]
- Georgios Kaissis, Marcus R. Makowski, Daniel Rückert, Rickmer Braren (2020) 'Secure, privacy-preserving and federated machine learning in medical imaging', doi:10.1038/s42256-020-0186-1. Available at: <https://openalex.org/W3033511014> [reputation: high] [citations: 1348] [access: open]

## Limitations

- *SupervisorAgent*: 171/150 words, 4/2 reputable sources; minimums met -> researching further

*Scope note: this is a deliberately small, local pipeline (free scholarly APIs + local retrieval + a local LLM). See the README for the design trade-offs.*

## Appendix: decision log

| # | agent          | kind       | reason                                                                                                                                                                                                                                                                           |
|---|----------------|------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | SummariseAgent | summarised | Summary and claims were derived exclusively from the provided abstract text, with page numbers set to 0 as only the abstract is available and no full text exists.                                                                                                               |
| 2 | SummariseAgent | summarised | The summary and claims were derived exclusively from the provided abstract text, with page 0 assigned to all claims since only the abstract is available.                                                                                                                        |
| 3 | SummariseAgent | summarised | The summary and claims were derived exclusively from the provided abstract text, which specifies that only the abstract is available (no full text). All claims are based on verbatim text from the abstract with page 0 assigned as instructed since no full text is available. |
| 4 | RankingAgent   | ranked     | 2007.06081v1 (arxiv.org): medium — Open-access preprint server; dominant venue for ML/physics/CS research but not peer-reviewed [registry]                                                                                                                                       |
| 5 | RankingAgent   | ranked     | 2207.02337v1 (arxiv.org): medium — Open-access preprint server; dominant venue for ML/physics/CS research but not peer-reviewed [registry]                                                                                                                                       |

|    |                 |            |                                                                                                                                                                                                                                                                                                                                                    |
|----|-----------------|------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 6  | RankingAgent    | ranked     | W2995022099 (Foundations and Trends® in Machine Learning): high — Foundations and Trends® in Machine Learning is a peer-reviewed journal series published by now publishers, with rigorous editorial standards and high impact in the machine learning field. [cache]                                                                              |
| 7  | RankingAgent    | ranked     | W3033511014 (Nature Machine Intelligence): high — Nature is a highly respected peer-reviewed journal, and Nature Machine Intelligence is a specific journal within the Nature publishing group known for high-quality research in machine learning and AI. [llm]                                                                                   |
| 8  | DiscoveryAgent  | discovered | the title and abstract highlight vertical federated learning, asynchronous execution, and privacy techniques, which are critical for identifying relevant search terms and broader research areas in distributed machine learning systems.                                                                                                         |
| 9  | DiscoveryAgent  | discovered | the title and abstract highlight the intersection of federated learning and transfer learning with a focus on security vulnerabilities and defense mechanisms, requiring keywords that capture specific security challenges and methods, and topics that encompass broader fields related to secure machine learning systems.                      |
| 10 | DiscoveryAgent  | discovered | mined W2995022099                                                                                                                                                                                                                                                                                                                                  |
| 11 | DiscoveryAgent  | discovered | The title highlights secure and privacy-preserving aspects of machine learning in medical imaging, indicating a focus on protecting sensitive health data while enabling collaborative learning across institutions. Keywords capture specific techniques and applications, while topics reflect broader fields where these concepts are relevant. |
| 12 | SupervisorAgent | retry      | 171/150 words, 4/2 reputable sources; minimums met -> researching further                                                                                                                                                                                                                                                                          |
| 13 | SupervisorAgent | converge   | 171/150 words, 4/2 reputable sources; minimums met; search exhausted (no new sources, no new topics); loop limit reached (1/1) -> publishing                                                                                                                                                                                                       |